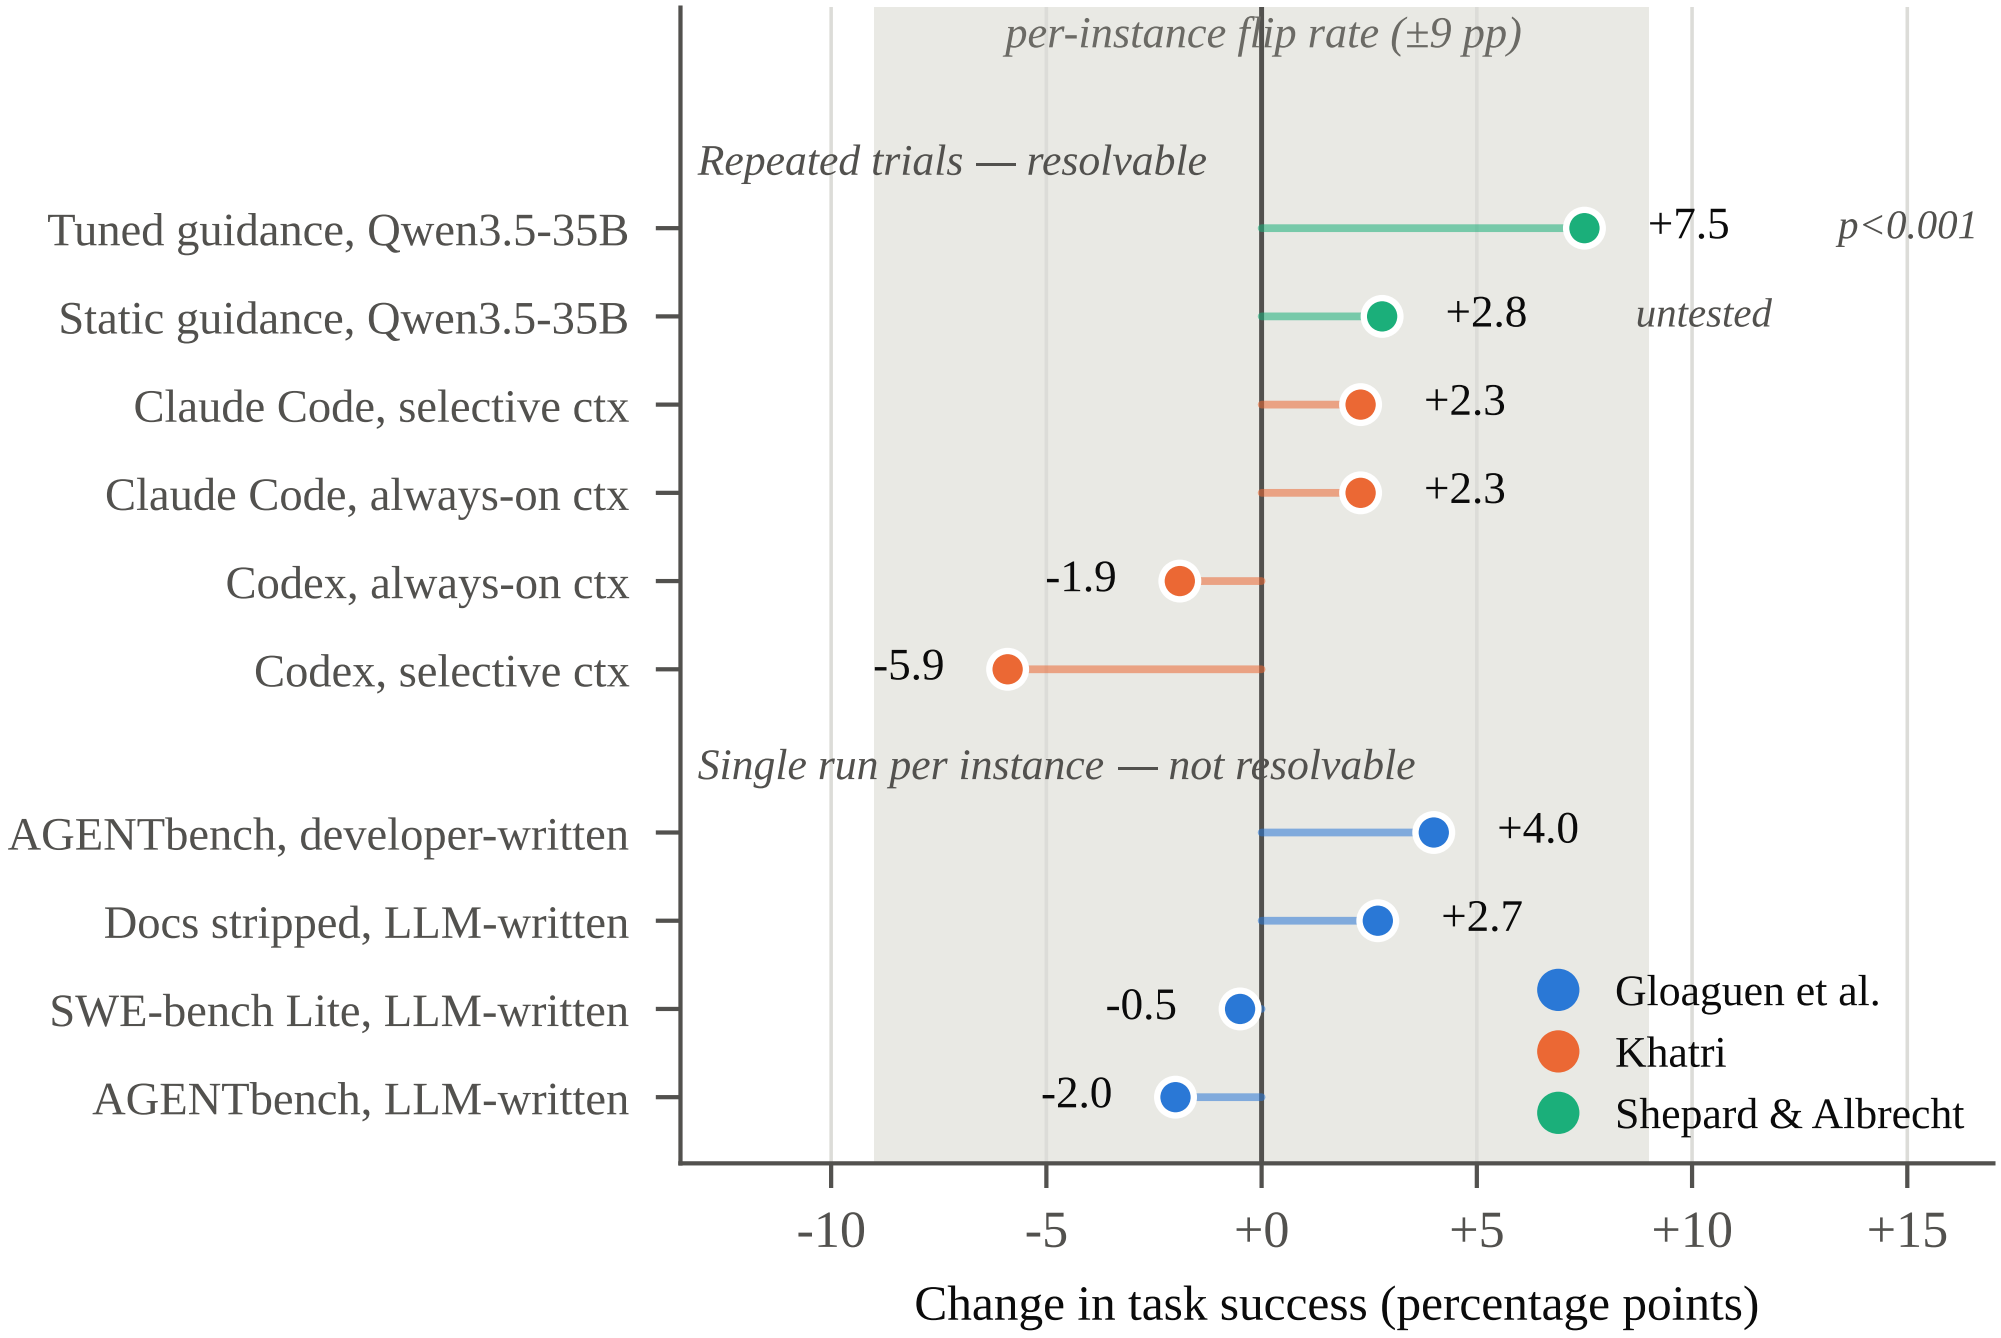
per-instance flip rate (±9 pp)
Repeated trials — resolvable
Tuned guidance, Qwen3.5-35B
+7.5
p<0.001
Static guidance, Qwen3.5-35B
+2.8
untested
Claude Code, selective ctx
+2.3
Claude Code, always-on ctx
+2.3
Codex, always-on ctx
-1.9
Codex, selective ctx
-5.9
Single run per instance — not resolvable
AGENTbench, developer-written
+4.0
Docs stripped, LLM-written
+2.7
SWE-bench Lite, LLM-written
-0.5
AGENTbench, LLM-written
-2.0
Gloaguen et al.
Khatri
Shepard & Albrecht
-10
-5
+0
+5
+10
+15
Change in task success (percentage points)