

EVIDENCE REVIEW · SEPTEMBER 2026

Does `/init` actually help?

What the research says about CLAUDE.md, AGENTS.md and every other repo context file your coding agent loads.

4

controlled ablations

15,549

agentic PRs analysed

247,694

instruction lifetimes

Your context file is two things. They have opposite signs.

Followed

Instructions work. Name a tool in the file and the agent uses it **1.6×** per task, versus under **0.01×** when unmentioned.

Ignored

Repository overviews do not. The architecture tour **/init** writes did not help agents find files any faster — and cost **20–23%** more inference.

Read the effect sizes with care: **~9% of outcomes flip between byte-identical runs** at temperature zero. Most published effects sit inside that noise.

Prose about code is a lossy picture of code.

4 / 45

Behavioural questions a natural-language summary of a file can answer. **The source itself answers 27 / 45.**

No fix

A frontier model's summary scores **exactly as poorly** as a 3B model's. The loss is in the representation — you cannot write your way out of it.

These files only ever grow.

+226%

Growth over a context file's lifetime — a net **+4.9 instructions per commit**, across 1,867 repositories.

Stuck

The older a line gets, the **less likely** anyone deletes it. Adding is cheap; proving a deletion is safe is not.

Catastrophic remembering.

Run it. Then delete most of it.

Keep

Only what the repo **cannot tell the agent itself**: non-obvious commands, gotchas, prohibitions, conventions that differ from the default.

Move

Enforcement into **hooks**. Procedures into **skills**. Module detail into **subdirectory files**. Then cut the architecture tour.

The practice is sound. **The default artefact is not.**