

# Does `/init` actually help?

What the research says about CLAUDE.md, AGENTS.md and every other repo context file your coding agent loads.

5

controlled ablations

15,549

agentic PRs analysed

247,694

instruction lifetimes

# Your context file is two things. They have opposite signs.

## Followed

**Instructions work.** Name a tool in the file and the agent uses it **1.6×** per task, versus under **0.01×** when unmentioned.

## Ignored

**Repository overviews do not.** The architecture tour `/init` writes did not help agents find files any faster — and cost **20-23%** more inference.

Read the effect sizes with care: **~9% of outcomes flip between byte-identical runs** at temperature zero — so most published numbers come from designs that cannot resolve the effect they report.

# Prose about code is a lossy picture of code.

4 / 45

Behavioural questions a natural-language summary of a file can answer. **The source itself answers 27 / 45.**

**No fix**

A frontier model's summary scores **exactly as poorly** as a 3B model's. The loss is in the representation — you cannot write your way out of it.

# These files only ever grow.

## +226%

Growth over a context file's lifetime — a net **+4.9 instructions per commit**, across 1,867 repositories.

## Stuck

The older a line gets, the **less likely** anyone deletes it. Adding is cheap; proving a deletion is safe is not.

**Catastrophic remembering.**

# One thing did beat the noise.

**25.5 → 33.0%**

Tasks resolved, after the guidance file was **tuned against the agent's own failures** — four trials,  **$p < 0.001$** . The untuned starting file gained **2.8 pts**, which the paper doesn't test.

## Coverage

The gain was **finding the right file** (+14.5 pts), not writing a better patch — patch quality never moved. And it was measured on an **open 35B model**, not a frontier agent.

So write down **what the agent got wrong**, not what your repository contains. The weaker your model, the more this matters.

# Run it. Then delete most of it.

## Keep

Only what the repo **cannot tell the agent itself**: non-obvious commands, gotchas, prohibitions, and every mistake you have had to correct twice.

## Move

Enforcement into **hooks**. Procedures into **skills**. Module detail into **subdirectory files**. Then cut the architecture tour.

The practice is sound. **The default artefact is not.**