

Tiel-Coder on the 36 GB box

A re-quantized Ornith-1.5 landed on the server, and its uncensored sibling with it. Both were measured against the field v3 recommended — on a runtime that has since changed the rules underneath all of it.

DEPLOY

Use `tiel-coder:35b-q5-ctx256k-agentic` — the variant, not the tag as it shipped. Same weights; the shipped tag carries `presence_penalty 1.5`, which costs 52% of generation speed and buys nothing measurable.


It is one of only two models that solved the *specification* rather than the visible tests — all 18 held-out tests, three runs from three. It holds 262,144 tokens at 34.13 GB, recalls at 254,181, and refuses an over-long prompt instead of quietly answering from half of it. Append `<|think_off|>` and it gets 2.3× faster for nothing.

111 tok/s generation Tiel, penalty removed	+52% over the tag as shipped one sampler setting	254,182 tokens of verified recall deepest in v1–v4
3/3 runs that solved the spec all 18 held-out tests	2.3× faster with thinking off no loss of correctness	7 models that silently halve an over-long prompt

- ✓ S1 TIEL
- ✓ S2 FIELD
- ✓ S3 SESSIONS
- ✓ S5 CYBERTIEL
- ✓ S6 SANDBOXED


Speed, and what it fails to predict

median of 3 · server idle before each



Time to finish the hard job

Median of three real Claude Code sessions on the three-module fixture. Lower is better.



Tokens per second does not predict this — the fastest model on the box is last.

What the presence penalty costs

Identical weights, one sampler setting apart.



Prefill moves less than 2%: the penalty is paid per generated token, in the sampler.

What thinking costs


Tiel on the hard fixture, median of three, with and without `<|think_off|>`.



Held-out scores were equal or better with thinking off, and the only session Tiel failed all day was a thinking-on run.

Cost of one agent turn

Seconds to first token when a long prefix is reused and only a tool result is appended.



New in Ollama 0.33.3. On 0.32.15 every one of these was a full re-read of the context.

THE FAILURE WITH NO ERROR MESSAGE

Past its context window, **7 of the older models silently keep half the prompt** ($\text{num_ctx}/2 + 2$ tokens) and answer anyway — ornith, north-mini, gemma4 and the qwen3.6 control among them. Both Tiel builds return **HTTP 400** instead. Claude Code cannot send `num_ctx`, so a model in the halving class will answer from half your repository with nothing in the transcript to say so.

Measured side by side

Both at Q5_K_XL with `presence_penalty 0` , same window, same harness.

	TIEL	CYBERTIEL
generation @2k	111.1 tok/s	110.0 tok/s
cold prefill @35k	3,483 tok/s	3,519 tok/s
deepest verified recall	254,181 tok	254,182 tok
resident @262,144	34.13 GB	34.13 GB
vision checks	25/25	25/25
over-long prompt	refuses (400)	refuses (400)
hard fixture, sandboxed	1/1 runs at 18/18	1/1 runs at 18/18
refusals on 8 benign security prompts	0	0

The answer: Tiel

Not because CyberTiel is worse — because it is indistinguishable, and costs more to run safely.

They measure the same. Generation within 1%, recall one token apart, identical memory, identical refusal of an over-long prompt, both 25/25 on vision. The abliteration and the cyber-weighted imatrix did not move anything this benchmark can see.

Neither refused any legitimate work. Across eight benign defensive-security prompts — SQL-injection fixes, password hashing, log triage, path-traversal validation — both answered all eight. The uncensored build buys nothing on the work you actually do.

But it costs more to host. An abliterated model has had its refusal behaviour removed, so its publisher tells you to sandbox it at the OS level. Here that meant a container with no host mounts, a read-only rootfs and an egress allowlist of exactly one address. That is real operational overhead for a model that measured the same. **Run Tiel. Reach for CyberTiel only if you hit a refusal that blocks legitimate work** — and then keep it in the sandbox.

MODEL		GEN TOK/S	COLD PREFILL	GATES	VISION	PAST ITS WINDOW
cascade-2	cut	141	7,017	8/10	no	halves silently
north-mini		136	4,592	10/10	no	halves silently
qwen3.6 35b		132	4,046	10/10	yes	halves silently
ornith 1.0		127	3,555	10/10	no	halves silently
nemotron-3.5-L		127	2,674	10/10	no	halves silently
Tiel · pp0		111	3,483	9/10	yes	refuses
CyberTiel · pp0		110	3,519	10/10	yes	refuses
gemma4 26b		109	6,126	10/10	yes	halves silently
Tiel · shipped		73	3,548	10/10	yes	refuses
qwen3.8 27b	cut	30	1,457	9/10	yes	halves silently

The visible tests are the ones the model can see. The held-out tests check the docstring specification it was asked to implement. A model can turn the first green while leaving the second failing — and most of them did.

MODEL	THINKING	HARNESS	PASSED	MEDIAN	HELD-OUT TESTS, PER RUN
CyberTiel · pp0	off	sandbox	1/1	37 s	18/18
ornith 1.0	on	host	3/3	47 s	18/18, 16/18, 17/18
Tiel · shipped	off	host	3/3	56 s	18/18, 18/18, 17/18
CyberTiel · pp0	on	sandbox	1/1	57 s	18/18
qwen3.6 35b	on	host	3/3	60 s	17/18, 15/18, 16/18
nemotron-3.5-L	on	host	3/3	74 s	14/18, 14/18, 13/18
Tiel · pp0	on	sandbox	1/1	80 s	18/18
Tiel · pp0	on	host	3/3	83 s	18/18, 18/18, 18/18
gemma4 26b	on	host	3/3	92 s	18/18, 18/18, 18/18
north-mini	on	host	3/3	126 s	18/18, 14/18, 17/18
Tiel · shipped	on	host	2/3	131 s	18/18, 16/18, 17/18
cascade-2	on	host	0/3	419 s	0/18, 3/18, 3/18
qwen3.8 27b	on	host	3/3	787 s	18/18, 18/18, 18/18

What we learned that we did not expect

The runtime started caching prompts, and a v3 rule died with it

Ollama 0.33.3 serves a repeated prefix from cache; 0.32.15 did not, and every v3 transcript reports `cache_read 0`. An agent turn now prefills only its new tail — 520 tokens instead of 30,042. v3 ranked this field on prefill because "the loop re-reads its context every turn". That is no longer true here.

Thinking is a 2.3× tax on tool-heavy work

Same model, same fixture: 131/148/117 s with thinking, 74/53/56 s without — and the held-out scores came back equal or better without it. The one session Tiel failed all day was a thinking-on run.

A clean gate battery is not a reliability number

Tiel passed all seven gates three times over. Re-running the nested-schema gate alone sixteen times gives 13/16. The battery samples once, at the tag's shipped temperature 0.6 — so every single-shot gate result in v1–v4 is one sample.

Abliteration changed almost nothing measurable

CyberTiel matches Tiel to within noise on every axis, and on eight benign defensive-security prompts *neither* model refused. The uncensored build earns its extra operational cost only if you actually hit a refusal.

The fastest model on the box is still the worst at the job

`cascade-2` leads both speed axes — 140.7 tok/s and 7,017 prefill — and scored 0 of 3 on the hard fixture, with held-out scores of 0, 3 and 3 out of 18, taking 419 s to do it. v3 rejected it for the same defect. Nothing changed, so it is cut.

Half this round's bugs were in the harness

A retrieval probe that queried `localhost` and scored three empty replies as model failures; a sandbox that never delivered its fixture, so every isolated session ran against an empty directory. Both produced plausible numbers instead of errors, which is what makes that class dangerous. All twelve are in `review.md`.

Every figure is generated from the TSVs in `results/` by `make-report.py`; method and caveats in `measurements.md`, the twelve harness defects found and fixed in `review.md`, exact model digests and versions in `results/provenance.txt`. Sessions are real `claude -p` runs scored from the repository afterwards, never from the model's own summary. Cut models stay in the tables and out of the headlines.